


Part 1/2



No Mercy for Manual Entry

29/Sep/2021
Workshop @ AMLD2021




Samurai



Wifi Info


Network: Free_STCC
User ID : 2317791957
Password: 5197



Samurai




Authors




Valerio Rossetti, PhD

Co-founder of SamurAI
Senior Data Scientist



Giulio Grossi, PhD

Senior Quantitative Portfolio Manager
at ONE swiss bank




SamurAI




Get to know the audience

- Knowledge in Supervised Learning and Computer Vision
- Your coding skills in Python
- Do you have practical applications of the techniques in this workshop?



Manual Entry

- A lot of information in paper/pdf documents:
invoices, contracts, personal information, surveys ...



Manual Entry

- A lot of information in paper/pdf documents: invoices, contracts, personal information, surveys ...
- These documents are then treated manually:
 - costly,
 - time-consuming and
 - error-prone

Manual image-to-data process

Philomed Healthcare Solutions
Chemin du Némorand 14
1208 Genève
info@philomed.com

WHITENING *Active*
HYPOALLERGENIC DENTIFRICES

To
Dentista Santa
81, Labarthe Ave
53011, Babel
info@dentista.com

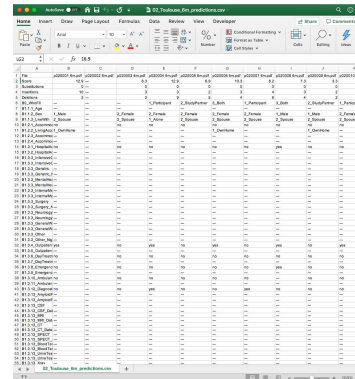
Facture
Facture No: 0854617
Date: 09/10/2020

Description	Code	Quantity	TVA	Price	Discount	Total
Levonorolol granule	20750	1	0	0	0	0
Gelsa F1 Chocout	C18	1	0	0	0	0
Gelsa F1 Chocout	C18	1	0	0	0	0

Contades Bancaires
Bancaires Philomed Healthcare Solutions Srl
BANQUE CHA 0284 7287 2087 201513
BIC:SWISS211 00000000000000000000
USD SWITZERLAND AG
SWITZERLAND AG, ZURICH

Total Hout: 16.00
Total TVA: 0.00
Total: 16.00
Discount: 0.00
Total en notre faveur: 16.00

Payé



Manual Entry

- A lot of information in paper/pdf documents: invoices, contracts, personal information, surveys ...
- These documents are then treated manually:
 - costly,
 - time-consuming and
 - error-prone
- Main solution: go fully digital and eliminate the paper document!

Manual image-to-data process

Philomed Healthcare Solutions
Chemin du Normandie 14
1208 Genève
info@philomed.com

WHITENING *Active*
HYPOALLERGENIC DENTIFRICES

To:
Dentista Slatia
81, Leclercq Ave
53011, Babel
info@dentista.com

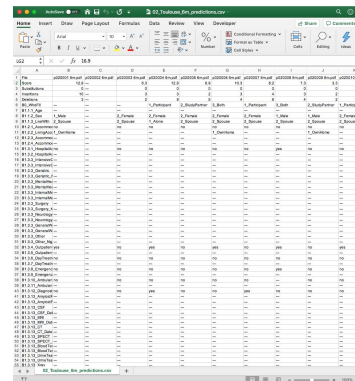
Facture
Facture No: 0854617
Date: 09/10/2020

Description	Code	Quantity	TVA	Price	Discount	Total
Levonorolol granule	20750	1	0	0	0	0
Gelsa F1 Choculat	C18	1	0	0	0	0
Gelsa F1 Choculat	C18	1	0	0	0	0

Controles Bancaires
Banque: Philomed Healthcare Solutions Slat
BANQUE CHIA 0284 7287 2287 245153
BIC: CHIA3331
USD SWITZERLAND AG
SWITZERLAND AG, ZURICH

Total Hors TVA: 16.00
Total TVA: 0.00
Total: 16.00
Discount: 0.00
Total en notre faveur: 16.00

Payé



Manual Entry: a zombie among us

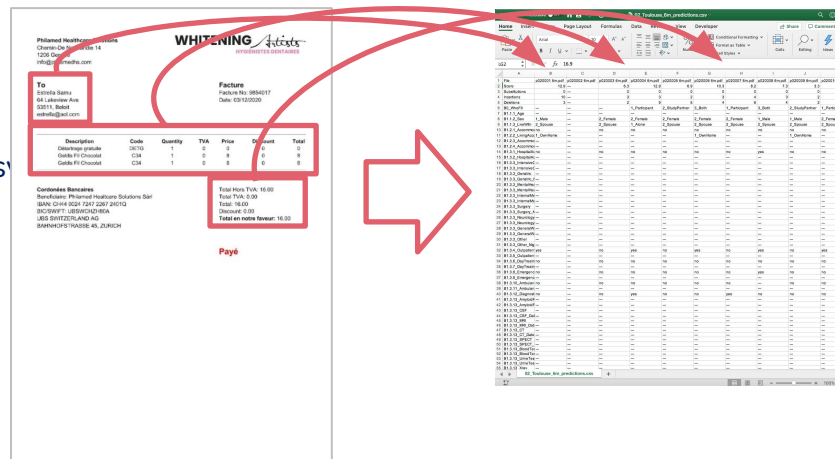
- In some cases we don't manage to eliminate the paper/pdf document
- Examples:
 - Traders send emails with pdf documents treated by the back-office
 - Commodity trading dealing with letters of credit
 - Banks dealing with client documents (passport, ID, scanned contract, ...)
 - Medical research dealing with huge volumes of paper documents from patients
 - Auditors oftentimes compare tables in pdf documents to Excel spreadsheets



This workshop: the goal

- ML can help you to automate processes with scanned paper documents
- This workshop presents few selected techniques to:
 - Classify documents
 - Extract information from documents
- Disclaimer:
 - This workshop doesn't present the fanciest / more powerful techniques to do these tasks
 - This workshop shows you few techniques that are easy to understand, implement and deploy as a beginner data scientist

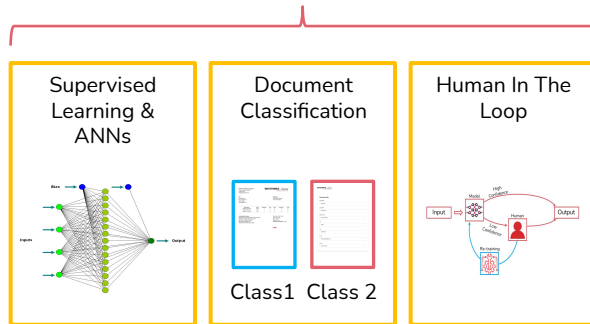
Automated image-to-data process



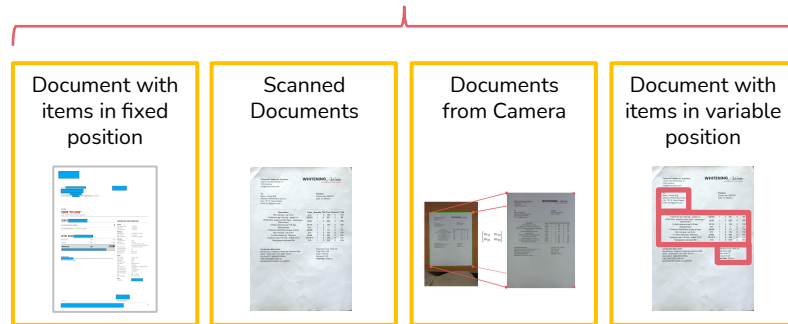
This workshop

- For beginners
- Get your hands dirty in code
- Adding levels of complexity


Part 1
(~1h30)



Part 2
(~2h00)



Workshop: Schedule



Session	Duration	Start - End	Subjects
Part 1	~1:30	13:30 – 15:00	Intro + Document Classification + HITL
Break	0:30	15:00 – 15:30	
Part 2	~2:00	15:30 – 17:30	Information Extraction

Your questions are welcomed!

Your 1st question



=

Swiss chocolate bar



«I don't get the big picture»

«I didn't understand what this code does»

«Why are we doing this?»



Supervised Learning

Attributes/
Features


Whitening Invoice

To: [Name], [Address], [City], [Country]

From: [Name], [Address], [City], [Country]

Description	Unit	Quantity	Unit Price	Amount	Total
[Description]	[Unit]	[Quantity]	[Unit Price]	[Amount]	[Total]
[Description]	[Unit]	[Quantity]	[Unit Price]	[Amount]	[Total]

Total: [Total Amount]



Predictions

It's a
document
Invoice!

Supervised Learning

Attributes/
Features

Personal Hardware Solutions
Computer Hardware & Software

WHITENING, JANEY

To: Mr. J. Smith
123 Main St
New York, NY 10001

From: Mr. J. Smith
123 Main St
New York, NY 10001

Subject: Personal Hardware Solutions

Dear Mr. Smith,

I am writing to you regarding the personal hardware solutions that we have for your business. We have a variety of options that can help you improve your productivity and efficiency.

Our solutions include:

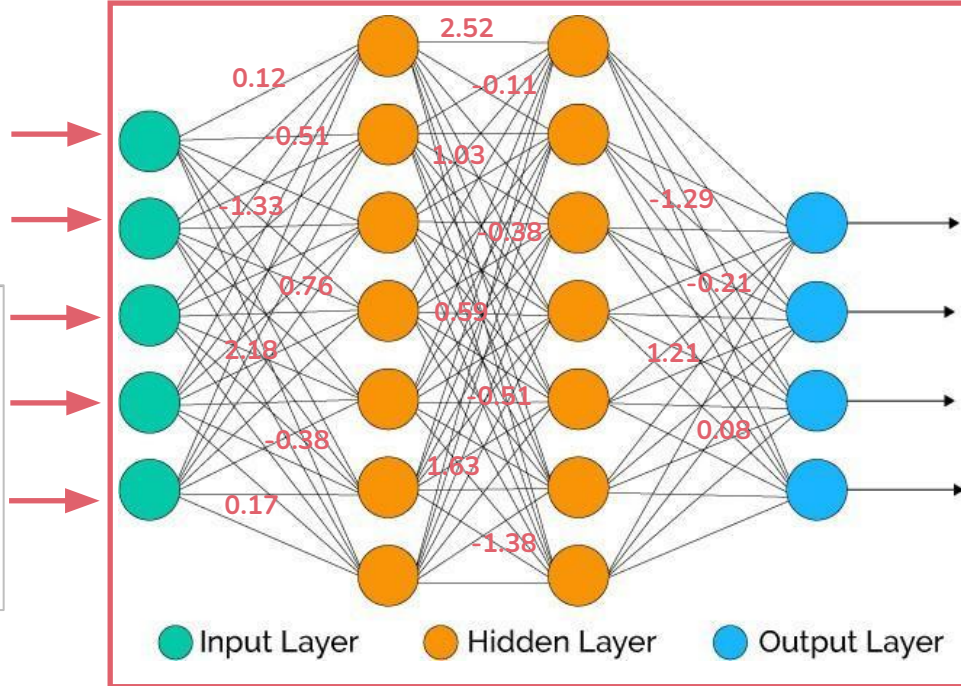
- Personalized hardware solutions
- Customized software solutions
- Cloud-based solutions
- Mobile solutions

We are confident that our solutions will meet your needs and provide you with the best possible results. Please contact us at 1-800-123-4567 for more information.

Sincerely,
Mr. J. Smith
123 Main St
New York, NY 10001

Personal Info

First Name: _____
Last Name: _____
Company: _____
Address: _____
City: _____
State: _____
Zip: _____
Phone: _____
Email: _____



Predictions

Labels

invoice
invoice
registration
invoice
registration
registration
invoice
invoice
registration

invoice
registration
invoice
registration
registration
invoice
registration
invoice
registration

Loss/CostFunction: 1000
Accuracy: 50%

Supervised Learning

Attributes/
Features

Personal Information

WHITENING, JANE

To: Mr. Jane White
123 Main St
New York, NY 10001

From: Mr. Jane White
123 Main St
New York, NY 10001

Subject: Personal Information

Personal Info

First Name:

Last Name:

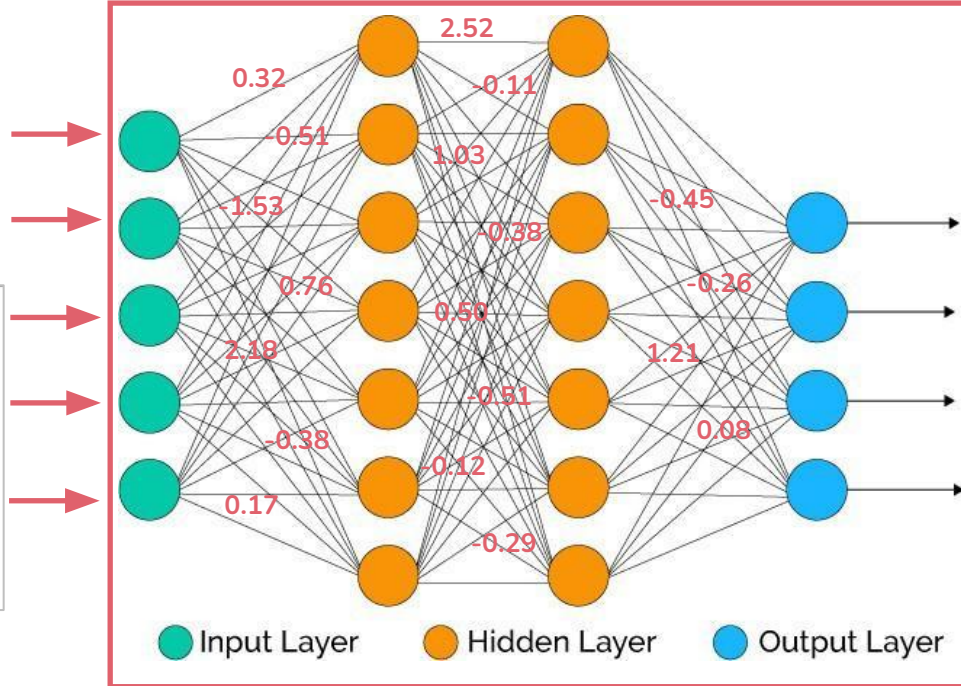
City:

State and Zip:

Country:

Phone:

Email:



Predictions

Labels

invoice
registration
registration
invoice
registration
invoice
invoice
invoice
registration

invoice
registration
invoice
registration
invoice
registration
invoice
invoice
registration

Loss/CostFunction: 600
Accuracy: 70%

Supervised Learning

Attributes/
Features

Personal Information

WHITENING, JANE

To: Mr. Jane Smith
123 Main St
New York, NY 10001

From: Mr. John Doe
456 Elm St
Los Angeles, CA 90001

Subject: Business Meeting

Date: 2023-10-27

Time: 10:00 AM

Location: Conference Room A

Personal Info

First Name:

Last Name:

Company:

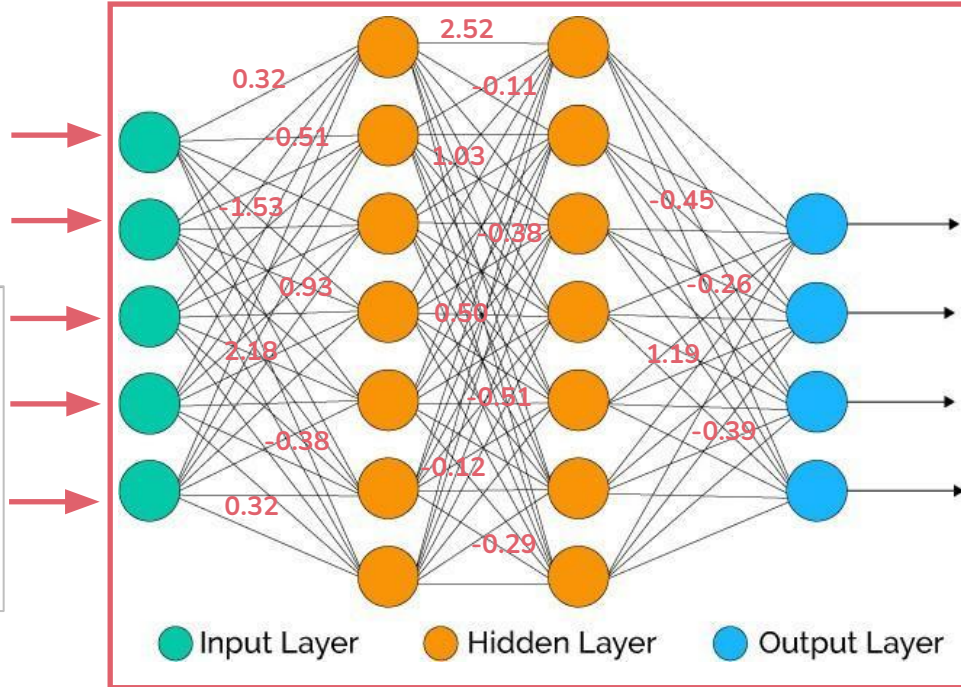
City:

State and Zip:

Country:

Phone:

Email:



Predictions

Labels

invoice
registration
invoice
registration
registration
invoice
invoice
invoice
registration



invoice
registration
invoice
registration
registration
invoice
registration
invoice
registration

Loss/CostFunction: 300
Accuracy: 90%

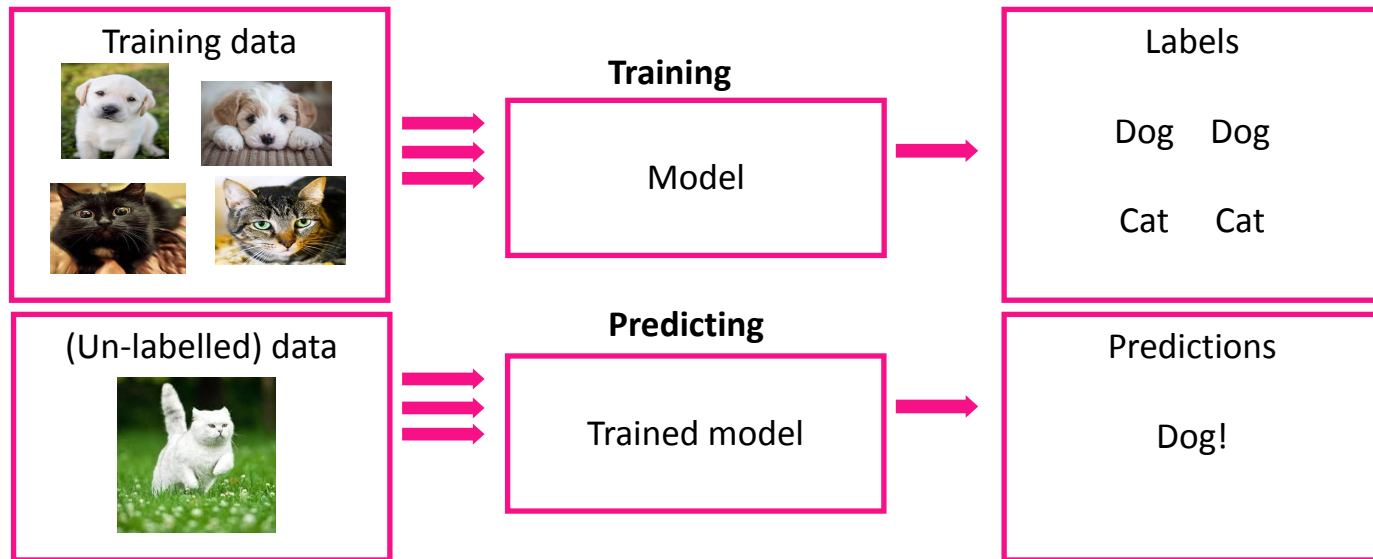
Training as a loss minimization

- The loss quantifies the spread between labels and predictions
- The optimizers are algorithms that find the (possibly absolute) minimum of the loss

During the training the optimizer finds the path to minimize the loss, like a river flowing downhill




Learning is hard



ANN: Dense Layers

- Multi-Layer Perceptrons are the simplest ANN
- Every node of a layer is connected with all the nodes in the previous and in the following layer
- These layers are called Fully connected or Dense





- Image of 500 x 500 x 3 pixels
- First layer 100 nodes
- Already for the first layer we need ~750k parameters

The MLP is great, but too many parameters!


ANN: Convolutional Layers

- We want to reduce the model parameters
- We introduce the concept of filter



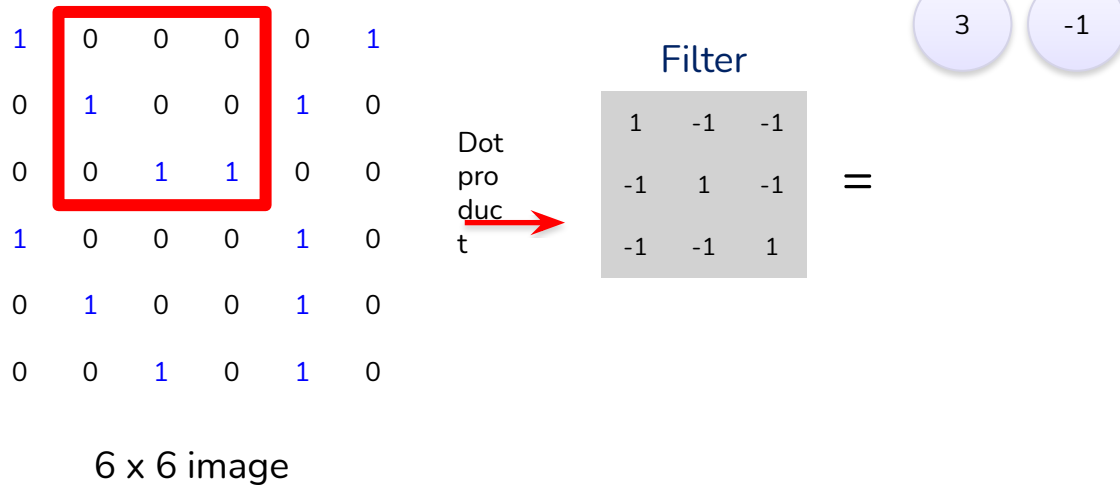
ANN: Convolutional Layers

- We want to reduce the model parameters
- We introduce the concept of filter



ANN: Convolutional Layers

- We want to reduce the model parameters
- We introduce the concept of filter



ANN: Convolutional Layers

- We want to reduce the model parameters
- We introduce the concept of filter

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

Dot
product
→

Filter


1	-1	-1
-1	1	-1
-1	-1	1

=

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

ANN: Convolutional Layers

- We want to reduce the model parameters
- We introduce the concept of filter



1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

Dot
product
→

Filter


1	-1	-1
-1	1	-1
-1	-1	1

=


3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

Dropout


- Regularization technique
- A fraction of the nodes are not considered in a training step
- This forces the network to have several “routes” in the nodes to ensure good performances



Standard Training



Training with Dropout




Dropout


- Regularization technique
- A fraction of the nodes are not considered in a training step
- This forces the network to have several “routes” in the nodes to ensure good performances

```
# Convolutional Layer 1
headmodel.add(Conv2D(8, (5, 5), padding='same', activation='relu'))
headmodel.add(Dropout(0.2))
```


This means “drop 20% of the nodes in the previous layer while training



Standard Training

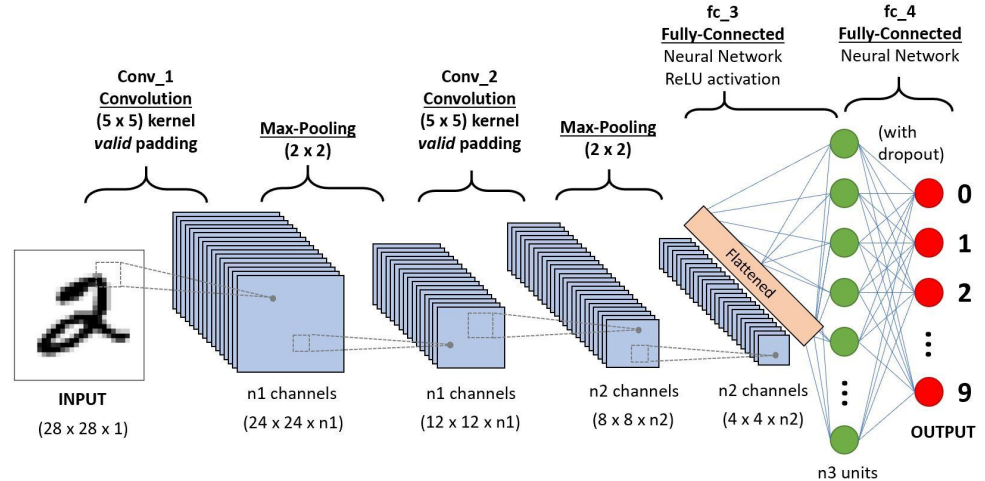


Training with Dropout



Typical CNN structure

- Typical CNNs are made of :
 - A series of convolutional layer + dropout + max-pooling
 - Fully-connected layers at the end
- Many parameters are subject to tuning:
 - Number and type of layer
 - Number of filters
 - Size of filters



CNN in Keras

```
def define_model(num_classes, epochs):  
    # Create the model  
    model = Sequential()  
  
    # Layer 1 (Convolutional)  
    model.add(Conv2D(4, (5, 5), input_shape=(X.shape[1], X.shape[2], 1), padding='same', activation='relu',  
                    model.add(Dropout(0.2))  
                    model.add(MaxPooling2D(pool_size=(2, 2)))  
  
    # Layer 2 (Convolutional)  
    model.add(Conv2D(4, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))  
    model.add(Dropout(0.2))  
    model.add(MaxPooling2D(pool_size=(2, 2)))  
  
    # Layer 3 (Convolutional)  
    #model.add(Conv2D(4, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))  
    #model.add(Dropout(0.2))  
    #model.add(MaxPooling2D(pool_size=(2, 2)))  
  
    # Additional Convolutional layers  
    # ...  
  
    # Additional Dense Layers  
    model.add(Flatten())  
    # model.add(Dense(6, activation='relu', kernel_constraint=maxnorm(3)))  
    model.add(Dense(num_classes, activation='softmax'))
```

Conv. Layer +
Dropout +
Max-Pooling

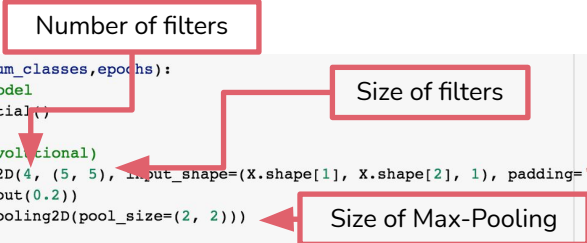
Conv. Layer +
Dropout +
Max-Pooling

Conv. Layer +
Dropout +
Max-Pooling

Fully Connected Layers

Output

CNN in Keras



```
def define_model(num_classes, epochs):  
    # Create the model  
    model = Sequential()  
  
    # Layer 1 (Convolutional)  
    model.add(Conv2D(4, (5, 5), input_shape=(X.shape[1], X.shape[2], 1), padding='same', activation='relu',  
                    # Number of filters  
                    # Size of filters  
                    model.add(Dropout(0.2))  
                    model.add(MaxPooling2D(pool_size=(2, 2))) # Size of Max-Pooling  
  
    # Layer 2 (Convolutional)  
    model.add(Conv2D(4, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))  
    model.add(Dropout(0.2))  
    model.add(MaxPooling2D(pool_size=(2, 2)))  
  
    # Layer 3 (Convolutional)  
    #model.add(Conv2D(4, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))  
    #model.add(Dropout(0.2))  
    #model.add(MaxPooling2D(pool_size=(2, 2)))  
  
    # Additional Convolutional layers  
    # ...  
  
    # Additional Dense Layers  
    model.add(Flatten())  
    # model.add(Dense(6, activation='relu', kernel_constraint=maxnorm(3)))  
    model.add(Dense(num_classes, activation='softmax'))
```

Document Classification

- We are going to train an algorithm to discriminate between different types of documents, only using their images
- The dataset:
 - ~400 invoices
 - ~400 registration
 - ~600 other
- We want the model to be robust enough to use images taken from a phone

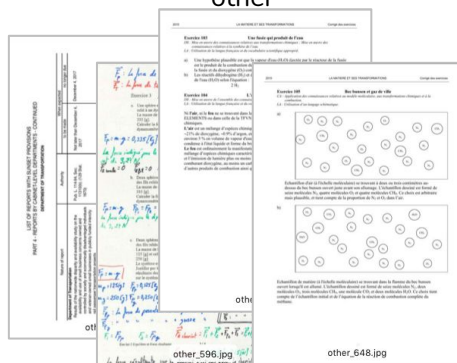
invoice

Invoice document from WHITENING. It includes a header with the company name and logo, followed by a table of items with columns for Description, Date, Quantity, Unit, Price, and Amount. Below the table is a section for 'Contenus Revenues' and a footer with contact information.

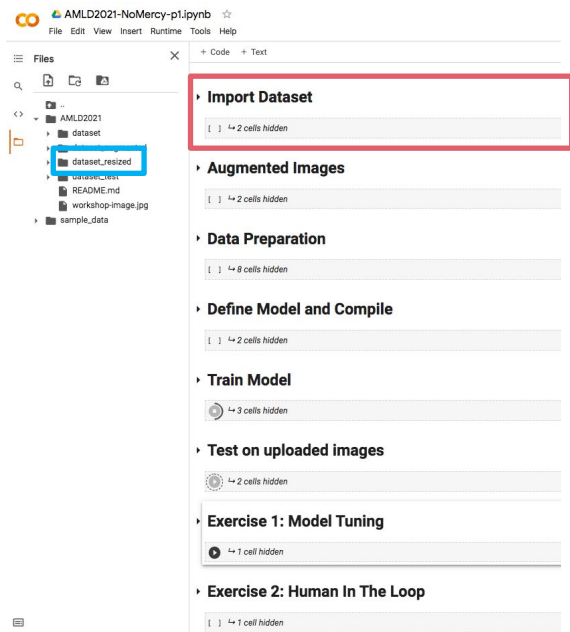
registration

Registration form from WHITENING. It contains fields for Personal info (First Name, Last Name, Date of Birth, Address, City, Country, Email, Phone) and a section for 'Choix de la couleur de la robe' (Choice of the color of the dress).

other

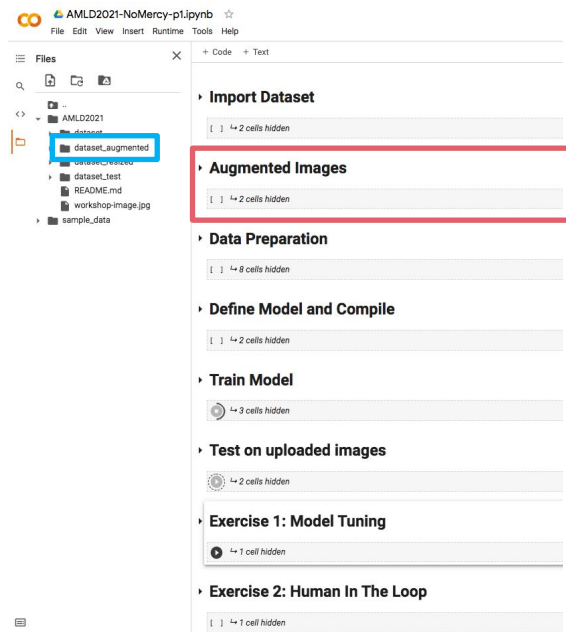


Notebook structure




- Import images from github
- On the left, you will see the folder AMLD2021 appear
- The main dataset used for this notebook is in [AMLD2021/dataset_resized/](#)
- ~1600images of 708x500 pixels

Notebook structure

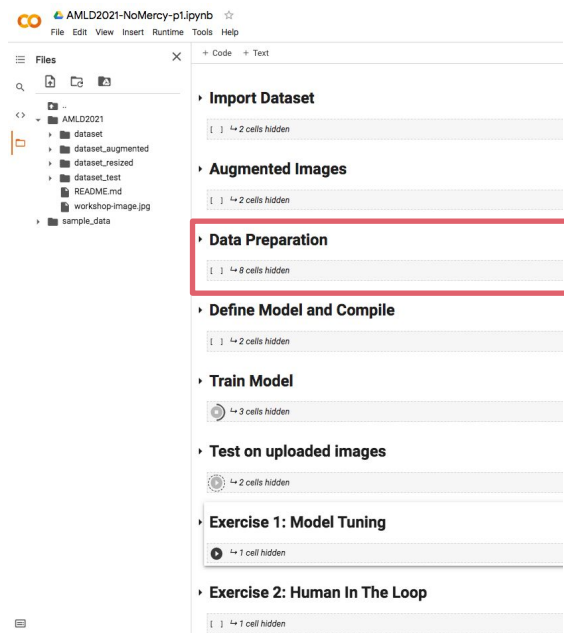


- More augmented images are created with distortions, brightness variations, tilt, ...



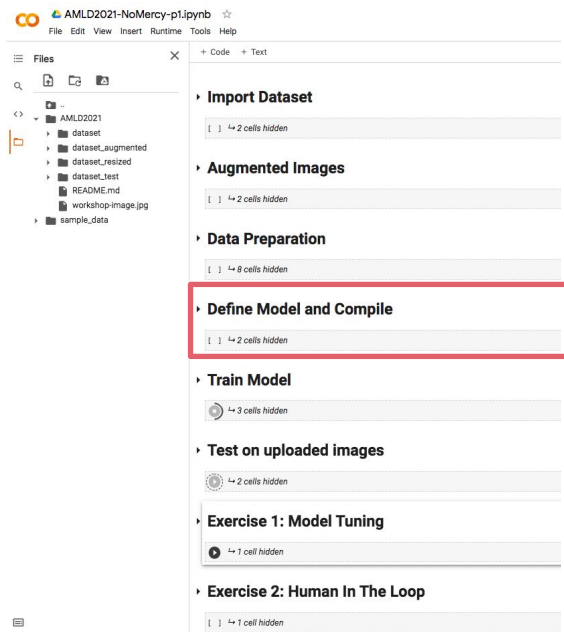
- Augmented images are moved to AMLD2021/dataset_augmented/

Notebook structure



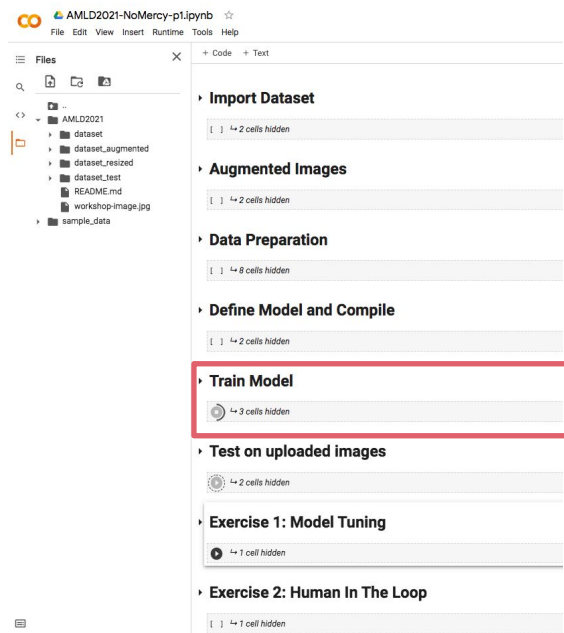
- Images are resized to a smaller size
- Dataset is created
- Train-test split

Notebook structure



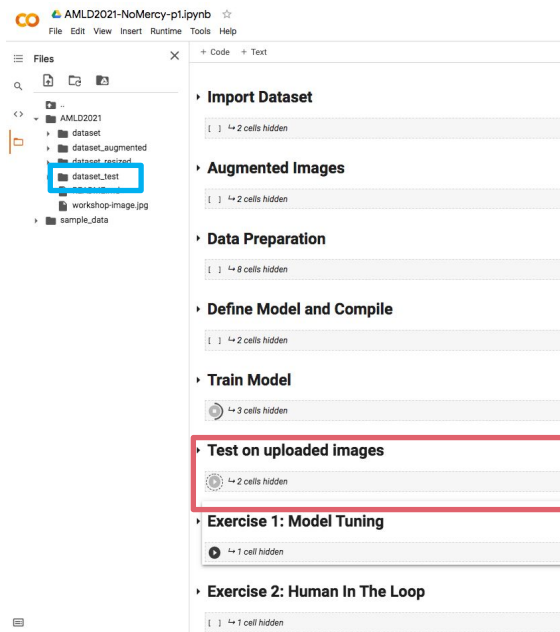
- Model is defined as a Keras ANN with Conv and Dense layers

Notebook structure

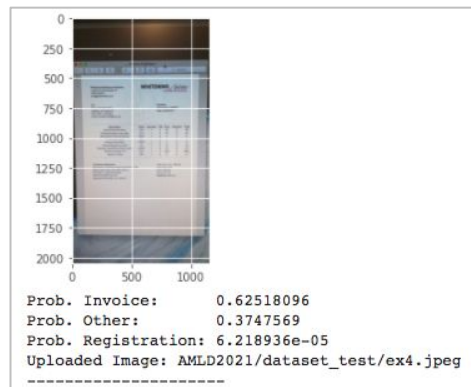


- Training and evaluation of performance

Notebook structure




- You can take a picture of a document and test the model on it
- Upload your picture to Colab and in the folder `AMLD2021/dataset_test/`
- The code will make predictions on all pictures in this folder




Exercise 1: Tuning the model

- Hyper-parameter tuning and feature engineering:
 - CNN layers, filters, dropout
 - Dense layers, number of nodes, dropout
 - Optimizer, Batch size, Num. of Epochs
 - Resize of the images (up in the data preparation)
- Goal: reach an accuracy of $> 99\%$ on the training and validation sets




Human in the loop (HITL)

- AI systems are typically <100% accurate
- We can keep a HITL to mitigate for mistakes
 - Manually do the tasks for which the model has low confidence (low h)
 -



Human in the loop (HITL)

- AI systems are typically <100% accurate
- We can keep a HITL to mitigate for mistakes
 - Manually do the tasks for which the model has low confidence (low h)
 - Prepare data to retrain the model




Exercise 2: Human in the loop

- Make photos of several documents (~5-10 per class)
 - Jpeg format is perfect
 - Advice: name the images according to their class (ex: my_invoice_1.jpg).
This will help for the rest of the exercise
- Upload them on Colab and copy them to the right folder for re-training (for example AMLD2021/dataset_resize/invoice/ for invoices)
 - example: `!cp my_invoice_* AMLD2021/dataset_resize/invoice/`
- Re-build augmented images, retrain the model, and re-evaluate performance

Possible Extensions: OCR

- In our classification we used only the images of our documents. We didn't use the text within the document
- Here an example of how to do document classification using only the text extracted with OCR:
[link](https://www.jmlr.org/papers/volume3/blei03a/blei03a.pdf)
 - Using an LDA model: <https://www.jmlr.org/papers/volume3/blei03a/blei03a.pdf>




Image


OCR
⇒

Philamed Healthcare Solutions
Chemin-De Normandie 14 1206
Genève info@philamedha.com To
Name: James Bull Address: 6640 N
Blue Gum St City: 70116, New Orleans
Email: jbull@gmail.com Facture
Facture No: 2954781 Date: 31/08/2021
Coordonnées Bancaires Bénéficiaire:
Philamed Healthcare Solutions Sàrl
IBAN: CH44 0024 7247 2297 2421Q
BIC: SWIFT UBSWCH2H80A UBS
SWITZERLAND AG
BAHNHOFSTRASSE 45, ZÜRICH
Total Hors TVA: 1818.70 Total TVA:
0.00 Total: 1818.70 Discount: 0.00
Total Net: 1818.70 Description Code
Quantity TVA Price Discount Total Rdy
manqué - par 5 min C15 1 0 159 0 159
Traitement par rhd dpl - enfant 0-6
DETE9 1 0 89 0 89 DEIN DEAL (payé
sur Dein Deal) - détartrage +
blanchiment DEIN 1 0 599 0 599
Forfait traitement par rhd dpl, DETR 1
0 159 0 159 Golda Brosse C31 1 0 7 0
7 Philamed Gold4White Brosse à dents
GAWB 1 0 7 9 0 7 9 Rdy manqué - par
5 min C15 1 0 159 0 159 In-office
bleaching - Diamond DDM 1 0 499 0
499 Traitement par rhd dpl - enfant
10-15 DETE15 1 0 120 0 120
Radiographie dentaire BW C18 1 0
18 8 0 18 8

Text



invoice



registration



-



OCR

Text

Doc2Vec

-0.7677
6.33419
0.00032
32.5298

Vectorized
representation

A diagram of a neural network structure. It consists of 10 nodes (black circles) arranged in a grid-like pattern. The nodes are connected by 15 magenta lines, representing weights. The connections form a complex web, including horizontal, vertical, and diagonal links, as well as some longer-range connections. Two input arrows point into the network from the left, and two output arrows point out to the right.

Classifier Model

invoice

Whitening 

PERSONAL INFORMATION:
 Name: _____
 Address: _____
 City: _____
 State: _____
 Zip: _____
 Phone: _____
 Fax: _____
 E-mail: _____

Product Information:
 Product Name: _____
 Product Code: _____
 Product Description: _____

Quantity:
 Unit: _____
 Quantity: _____

Price:
 Unit Price: _____
 Total Price: _____

Comments:
 Comments: _____
 Comments: _____
 Comments: _____
 Comments: _____
 Comments: _____

Signature:
 Signature: _____
 Signature: _____
 Signature: _____
 Signature: _____
 Signature: _____

Printed:
 Printed: _____
 Printed: _____
 Printed: _____
 Printed: _____
 Printed: _____

registration

[illegible]